

ML Pipeline Preparation

July 14, 2020

1 ML Pipeline Preparation

Follow the instructions below to help you create your ML pipeline. ### 1. Import libraries and load data from database. - Import Python libraries - Load dataset from database with `read_sql_table` - Define feature and target variables X and Y

```
In [1]: # import libraries
import nltk
nltk.download('punkt')
nltk.download('wordnet')
nltk.download('averaged_perceptron_tagger')

import plotly

import numpy as np
import pandas as pd
pd.set_option('display.max_columns', 500)

import sys
import os
import re
from sqlalchemy import create_engine
import pickle

from scipy.stats import gmean
from sklearn.pipeline import Pipeline, FeatureUnion
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import classification_report, confusion_matrix, fbeta_score, make_s
from sklearn.ensemble import GradientBoostingClassifier, RandomForestClassifier, AdaBoos
from sklearn.feature_extraction.text import TfidfTransformer, CountVectorizer
from sklearn.multioutput import MultiOutputClassifier
from sklearn.base import BaseEstimator, TransformerMixin

[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Unzipping tokenizers/punkt.zip.
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Unzipping corpora/wordnet.zip.
```

```
[nltk_data] Downloading package averaged_perceptron_tagger to
[nltk_data] /root/nltk_data...
[nltk_data] Unzipping taggers/averaged_perceptron_tagger.zip.
```

```
In [2]: # load data from database
        database_filepath = "./disaster_response_db.db"
        engine = create_engine('sqlite:/// ' + database_filepath)
        table_name = os.path.basename(database_filepath).replace(".db","") + "_table"
        df = pd.read_sql_table(table_name,engine)
```

```
In [3]: # 'child_alone' = Unary column = all Zero values
        # 'related' should be 0,1 but contains 2
        df.describe()
```

```
Out[3]:
```

	id	related	request	offer	aid_related \
count	26216.00000	26216.000000	26216.000000	26216.000000	26216.000000
mean	15224.82133	0.773650	0.170659	0.004501	0.414251
std	8826.88914	0.435276	0.376218	0.066940	0.492602
min	2.00000	0.000000	0.000000	0.000000	0.000000
25%	7446.75000	1.000000	0.000000	0.000000	0.000000
50%	15662.50000	1.000000	0.000000	0.000000	0.000000
75%	22924.25000	1.000000	0.000000	0.000000	1.000000
max	30265.00000	2.000000	1.000000	1.000000	1.000000

	medical_help	medical_products	search_and_rescue	security \
count	26216.000000	26216.000000	26216.000000	26216.000000
mean	0.079493	0.050084	0.027617	0.017966
std	0.270513	0.218122	0.163875	0.132831
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000

	military	child_alone	water	food	shelter \
count	26216.000000	26216.0	26216.000000	26216.000000	26216.000000
mean	0.032804	0.0	0.063778	0.111497	0.088267
std	0.178128	0.0	0.244361	0.314752	0.283688
min	0.000000	0.0	0.000000	0.000000	0.000000
25%	0.000000	0.0	0.000000	0.000000	0.000000
50%	0.000000	0.0	0.000000	0.000000	0.000000
75%	0.000000	0.0	0.000000	0.000000	0.000000
max	1.000000	0.0	1.000000	1.000000	1.000000

	clothing	money	missing_people	refugees	death \
count	26216.000000	26216.000000	26216.000000	26216.000000	26216.000000
mean	0.015449	0.023039	0.011367	0.033377	0.045545

std	0.123331	0.150031	0.106011	0.179621	0.208500
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

	other_aid	infrastructure_related	transport	buildings	\
count	26216.000000	26216.000000	26216.000000	26216.000000	
mean	0.131446	0.065037	0.045812	0.050847	
std	0.337894	0.246595	0.209081	0.219689	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	

	electricity	tools	hospitals	shops	aid_centers	\
count	26216.000000	26216.000000	26216.000000	26216.000000	26216.000000	
mean	0.020293	0.006065	0.010795	0.004577	0.011787	
std	0.141003	0.077643	0.103338	0.067502	0.107927	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	1.000000	

	other_infrastructure	weather_related	floods	storm	\
count	26216.000000	26216.000000	26216.000000	26216.000000	
mean	0.043904	0.278341	0.082202	0.093187	
std	0.204887	0.448191	0.274677	0.290700	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	1.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	

	fire	earthquake	cold	other_weather	direct_report
count	26216.000000	26216.000000	26216.000000	26216.000000	26216.000000
mean	0.010757	0.093645	0.020217	0.052487	0.193584
std	0.103158	0.291340	0.140743	0.223011	0.395114
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

In [4]: # 'related' = 2 = 188 rows = assume to be anomaly and remove all

```
df.groupby("related").count()
```

```
Out[4]:
```

	id	message	original	genre	request	offer	aid_related	\
related								
0	6122	6122	3395	6122	6122	6122	6122	
1	19906	19906	6643	19906	19906	19906	19906	
2	188	188	132	188	188	188	188	

	medical_help	medical_products	search_and_rescue	security	\
related					
0	6122	6122	6122	6122	6122
1	19906	19906	19906	19906	19906
2	188	188	188	188	188

	military	child_alone	water	food	shelter	clothing	money	\
related								
0	6122	6122	6122	6122	6122	6122	6122	
1	19906	19906	19906	19906	19906	19906	19906	
2	188	188	188	188	188	188	188	

	missing_people	refugees	death	other_aid	infrastructure_related	\
related						
0	6122	6122	6122	6122	6122	6122
1	19906	19906	19906	19906	19906	19906
2	188	188	188	188	188	188

	transport	buildings	electricity	tools	hospitals	shops	\
related							
0	6122	6122	6122	6122	6122	6122	
1	19906	19906	19906	19906	19906	19906	
2	188	188	188	188	188	188	

	aid_centers	other_infrastructure	weather_related	floods	storm	\
related						
0	6122	6122	6122	6122	6122	6122
1	19906	19906	19906	19906	19906	19906
2	188	188	188	188	188	188

	fire	earthquake	cold	other_weather	direct_report
related					
0	6122	6122	6122	6122	6122
1	19906	19906	19906	19906	19906
2	188	188	188	188	188

```
In [5]: # Dropped all 188 rows
df = df[df.related != 2]
df.groupby("related").count()
```

```
Out[5]:
```

	id	message	original	genre	request	offer	aid_related	\
--	----	---------	----------	-------	---------	-------	-------------	---

related							
0	6122	6122	3395	6122	6122	6122	6122
1	19906	19906	6643	19906	19906	19906	19906

medical_help medical_products search_and_rescue security \							
related							
0	6122	6122	6122	6122	6122	6122	
1	19906	19906	19906	19906	19906	19906	

military child_alone water food shelter clothing money \							
related							
0	6122	6122	6122	6122	6122	6122	6122
1	19906	19906	19906	19906	19906	19906	19906

missing_people refugees death other_aid infrastructure_related \							
related							
0	6122	6122	6122	6122	6122	6122	6122
1	19906	19906	19906	19906	19906	19906	19906

transport buildings electricity tools hospitals shops \							
related							
0	6122	6122	6122	6122	6122	6122	6122
1	19906	19906	19906	19906	19906	19906	19906

aid_centers other_infrastructure weather_related floods storm \							
related							
0	6122	6122	6122	6122	6122	6122	6122
1	19906	19906	19906	19906	19906	19906	19906

fire earthquake cold other_weather direct_report							
related							
0	6122	6122	6122	6122	6122	6122	6122
1	19906	19906	19906	19906	19906	19906	19906

```
In [6]: # Extract X and y variables from the data for the modelling
X = df['message']
y = df.iloc[:,4:]
```

1.0.1 2. Write a tokenization function to process your text data

```
In [7]: def tokenize(text,url_place_holder_string="urlplaceholder"):

    # Replace urls with url placeholder string
    url_regex = 'http[s]?://(?:[a-zA-Z]|[0-9]|[$-_@.&+]|[*\(\)\,]|(?:%[0-9a-fA-F][0-9a-f

    # Extract urls from the provided text
    detected_urls = re.findall(url_regex, text)
```

```

    # Replace url with url placeholder string
    for detected_url in detected_urls:
        text = text.replace(detected_url, url_place_holder_string)

    # Extract the word tokens from the provided text
    tokens = nltk.word_tokenize(text)

    # Lemmatizer to remove inflectional and derivationally related forms of a word
    lemmatizer = nltk.WordNetLemmatizer()

    # List of clean tokens
    clean_tokens = [lemmatizer.lemmatize(w).lower().strip() for w in tokens]

    return clean_tokens

In [8]: # Custom transformer to extract the starting verb of sentence
class StartingVerbExtractor(BaseEstimator, TransformerMixin):
    """
    Starting Verb Extractor class

    Extract first verb of sentence to convert it into new feature for the ML classifier
    """

    def starting_verb(self, text):
        sentence_list = nltk.sent_tokenize(text)
        for sentence in sentence_list:
            pos_tags = nltk.pos_tag(tokenize(sentence))
            first_word, first_tag = pos_tags[0]
            if first_tag in ['VB', 'VBP'] or first_word == 'RT':
                return True
        return False

    def fit(self, X, y=None):
        return self

    def transform(self, X):
        X_tagged = pd.Series(X).apply(self.starting_verb)
        return pd.DataFrame(X_tagged)

```

1.0.2 3. Build a machine learning pipeline

This machine pipeline should take in the message column as input and output classification results on the other 36 categories in the dataset. You may find the [MultiOutputClassifier](#) helpful for predicting multiple target variables.

```

In [9]: pipeline1 = Pipeline([
        ('features', FeatureUnion([

```

```

        ('text_pipeline', Pipeline([
            ('count_vectorizer', CountVectorizer(tokenizer=tokenize)),
            ('tfidf_transformer', TfidfTransformer())
        ]))
    ])),

    ('classifier', MultiOutputClassifier(AdaBoostClassifier()))
])

pipeline2 = Pipeline([
    ('features', FeatureUnion([

        ('text_pipeline', Pipeline([
            ('count_vectorizer', CountVectorizer(tokenizer=tokenize)),
            ('tfidf_transformer', TfidfTransformer())
        ])),

        ('starting_verb_transformer', StartingVerbExtractor())
    ])),

    ('classifier', MultiOutputClassifier(AdaBoostClassifier()))
])

```

1.0.3 4. Train pipeline

- Split data into train and test sets
- Train pipeline

```

In [10]: X_train, X_test, y_train, y_test = train_test_split(X, y)
        pipeline_fitted = pipeline1.fit(X_train, y_train)

```

1.0.4 5. Test your model

Report the f1 score, precision and recall for each output category of the dataset. You can do this by iterating through the columns and calling sklearn's `classification_report` on each.

```

In [11]: y_prediction_train = pipeline_fitted.predict(X_train)
        y_prediction_test = pipeline_fitted.predict(X_test)

        # Print classification report on test data
        print(classification_report(y_test.values, y_prediction_test, target_names=y.columns.values))

```

	precision	recall	f1-score	support
related	0.83	0.94	0.88	4982
request	0.75	0.53	0.62	1099
offer	0.00	0.00	0.00	34
aid_related	0.77	0.61	0.68	2688
medical_help	0.59	0.27	0.37	516

medical_products	0.73	0.36	0.48	329
search_and_rescue	0.50	0.17	0.25	172
security	0.21	0.04	0.07	120
military	0.58	0.36	0.44	219
child_alone	0.00	0.00	0.00	0
water	0.76	0.67	0.71	381
food	0.80	0.68	0.73	715
shelter	0.78	0.55	0.65	594
clothing	0.77	0.46	0.58	102
money	0.55	0.36	0.44	163
missing_people	0.67	0.25	0.36	64
refugees	0.55	0.22	0.32	224
death	0.71	0.46	0.56	297
other_aid	0.51	0.17	0.26	848
infrastructure_related	0.46	0.07	0.12	455
transport	0.57	0.17	0.27	288
buildings	0.64	0.39	0.48	308
electricity	0.45	0.19	0.26	135
tools	0.00	0.00	0.00	43
hospitals	0.33	0.07	0.12	86
shops	0.00	0.00	0.00	28
aid_centers	0.32	0.07	0.11	86
other_infrastructure	0.33	0.09	0.14	293
weather_related	0.85	0.66	0.74	1805
floods	0.88	0.58	0.70	567
storm	0.74	0.54	0.62	597
fire	0.53	0.32	0.40	60
earthquake	0.89	0.78	0.83	602
cold	0.77	0.32	0.45	136
other_weather	0.44	0.15	0.22	327
direct_report	0.70	0.48	0.57	1233
avg / total	0.73	0.59	0.63	20596

```
/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWarning
```

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.

```
/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1137: UndefinedMetricWarning
```

Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples.

```
In [12]: # Print classification report on training data
         print('\n',classification_report(y_train.values, y_prediction_train, target_names=y.col
```


	precision	recall	f1-score	support
related	0.83	0.95	0.88	14924
request	0.80	0.56	0.66	3375
offer	0.40	0.10	0.15	84
aid_related	0.77	0.61	0.68	8172
medical_help	0.66	0.29	0.40	1568
medical_products	0.70	0.35	0.47	984
search_and_rescue	0.64	0.19	0.29	552
security	0.47	0.08	0.13	351
military	0.66	0.38	0.48	641
child_alone	0.00	0.00	0.00	0
water	0.78	0.66	0.72	1291
food	0.82	0.69	0.75	2208
shelter	0.80	0.57	0.67	1720
clothing	0.74	0.46	0.57	303
money	0.65	0.34	0.44	441
missing_people	0.75	0.19	0.31	234
refugees	0.67	0.27	0.39	651
death	0.78	0.47	0.59	897
other_aid	0.58	0.18	0.28	2598
infrastructure_related	0.50	0.09	0.15	1250
transport	0.77	0.27	0.40	913
buildings	0.71	0.42	0.53	1025
electricity	0.72	0.31	0.43	397
tools	0.62	0.09	0.15	116
hospitals	0.44	0.11	0.18	197
shops	0.53	0.09	0.15	92
aid_centers	0.43	0.09	0.14	223
other_infrastructure	0.43	0.11	0.18	858
weather_related	0.87	0.65	0.75	5492
floods	0.88	0.55	0.68	1588
storm	0.76	0.56	0.65	1846
fire	0.74	0.33	0.45	222
earthquake	0.90	0.78	0.83	1853
cold	0.76	0.34	0.47	394
other_weather	0.57	0.19	0.28	1049
direct_report	0.76	0.50	0.60	3842
avg / total	0.77	0.60	0.65	62351

/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWarning

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.

/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1137: UndefinedMetricWarning

Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples.

1.0.5 6. Improve your model

Use grid search to find better parameters.

```
In [13]: # pipeline1.get_params().keys()
parameters_grid = {'classifier__estimator__learning_rate': [0.01, 0.02, 0.05],
                   'classifier__estimator__n_estimators': [10, 20, 40]}

#parameters_grid = {'classifier__estimator__n_estimators': [10, 20, 40]}

cv = GridSearchCV(pipeline1, param_grid=parameters_grid, scoring='f1_micro', n_jobs=-1)
cv.fit(X_train, y_train)

Out[13]: GridSearchCV(cv=None, error_score='raise',
                      estimator=Pipeline(memory=None,
                      steps=[('features', FeatureUnion(n_jobs=1,
                      transformer_list=[('text_pipeline', Pipeline(memory=None,
                      steps=[('count_vectorizer', CountVectorizer(analyzer='word', binary=False, decode_
                      dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',
                      lowercase...mator=None,
                      learning_rate=1.0, n_estimators=50, random_state=None),
                      n_jobs=1))]),
                      fit_params=None, iid=True, n_jobs=-1,
                      param_grid={'classifier__estimator__learning_rate': [0.01, 0.02, 0.05], 'classif
                      pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                      scoring='f1_micro', verbose=0)

In [14]: # Get the prediction values from the grid search cross validator
y_prediction_test = cv.predict(X_test)
y_prediction_train = cv.predict(X_train)
```

1.0.6 7. Test your model

Show the accuracy, precision, and recall of the tuned model.

Since this project focuses on code quality, process, and pipelines, there is no minimum performance metric needed to pass. However, make sure to fine tune your models for accuracy, precision and recall to make your project stand out - especially for your portfolio!

```
In [15]: # Print classification report on test data
print(classification_report(y_test.values, y_prediction_test, target_names=y.columns.va

precision    recall  f1-score   support
```

related	0.77	1.00	0.87	4982
request	0.58	0.38	0.46	1099
offer	0.00	0.00	0.00	34
aid_related	0.81	0.19	0.31	2688
medical_help	0.63	0.13	0.21	516
medical_products	0.72	0.13	0.22	329
search_and_rescue	0.60	0.19	0.29	172
security	0.00	0.00	0.00	120
military	0.51	0.15	0.23	219
child_alone	0.00	0.00	0.00	0
water	0.56	0.82	0.67	381
food	0.77	0.68	0.72	715
shelter	0.86	0.31	0.46	594
clothing	0.74	0.33	0.46	102
money	0.50	0.12	0.20	163
missing_people	0.67	0.31	0.43	64
refugees	0.50	0.01	0.03	224
death	0.78	0.17	0.28	297
other_aid	0.00	0.00	0.00	848
infrastructure_related	0.00	0.00	0.00	455
transport	0.53	0.19	0.28	288
buildings	0.00	0.00	0.00	308
electricity	0.00	0.00	0.00	135
tools	0.00	0.00	0.00	43
hospitals	0.00	0.00	0.00	86
shops	0.00	0.00	0.00	28
aid_centers	0.00	0.00	0.00	86
other_infrastructure	0.00	0.00	0.00	293
weather_related	0.91	0.22	0.35	1805
floods	0.94	0.36	0.52	567
storm	0.77	0.28	0.41	597
fire	0.44	0.40	0.42	60
earthquake	0.90	0.65	0.75	602
cold	0.82	0.26	0.40	136
other_weather	0.59	0.11	0.19	327
direct_report	0.60	0.39	0.47	1233
avg / total	0.66	0.44	0.47	20596

/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWarning

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.

/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1137: UndefinedMetricWarning

Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples.

```
In [16]: # Print classification report on training data
print('\n',classification_report(y_train.values, y_prediction_train, target_names=y.col
```

	precision	recall	f1-score	support
related	0.76	1.00	0.87	14924
request	0.62	0.38	0.47	3375
offer	1.00	0.02	0.05	84
aid_related	0.80	0.19	0.31	8172
medical_help	0.64	0.11	0.19	1568
medical_products	0.73	0.13	0.21	984
search_and_rescue	0.65	0.19	0.29	552
security	0.00	0.00	0.00	351
military	0.52	0.14	0.22	641
child_alone	0.00	0.00	0.00	0
water	0.58	0.85	0.69	1291
food	0.78	0.68	0.72	2208
shelter	0.84	0.29	0.43	1720
clothing	0.75	0.31	0.44	303
money	0.57	0.17	0.27	441
missing_people	0.68	0.24	0.36	234
refugees	0.83	0.02	0.03	651
death	0.74	0.16	0.27	897
other_aid	0.00	0.00	0.00	2598
infrastructure_related	0.00	0.00	0.00	1250
transport	0.59	0.26	0.37	913
buildings	0.00	0.00	0.00	1025
electricity	0.00	0.00	0.00	397
tools	0.00	0.00	0.00	116
hospitals	0.00	0.00	0.00	197
shops	0.00	0.00	0.00	92
aid_centers	0.00	0.00	0.00	223
other_infrastructure	0.00	0.00	0.00	858
weather_related	0.92	0.22	0.36	5492
floods	0.92	0.33	0.48	1588
storm	0.72	0.25	0.38	1846
fire	0.53	0.41	0.46	222
earthquake	0.90	0.64	0.75	1853
cold	0.64	0.20	0.31	394
other_weather	0.57	0.12	0.21	1049
direct_report	0.64	0.39	0.48	3842
avg / total	0.67	0.44	0.47	62351

/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWarning

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.

/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1137: UndefinedMetricWarning

Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples.

1.0.7 8. Try improving your model further. Here are a few ideas:

- try other machine learning algorithms
- add other features besides the TF-IDF

```
In [17]: #Use pipeline2 which includes StartingVerbEstimator
X_train, X_test, y_train, y_test = train_test_split(X, y)
pipeline_fitted = pipeline2.fit(X_train, y_train)

y_prediction_train = pipeline_fitted.predict(X_train)
y_prediction_test = pipeline_fitted.predict(X_test)

# Print classification report on test data
print(classification_report(y_test.values, y_prediction_test, target_names=y.columns.values))
```

	precision	recall	f1-score	support
related	0.82	0.95	0.88	4938
request	0.77	0.53	0.63	1119
offer	0.14	0.04	0.06	24
aid_related	0.76	0.61	0.68	2682
medical_help	0.54	0.25	0.35	502
medical_products	0.60	0.33	0.43	311
search_and_rescue	0.69	0.21	0.32	181
security	0.29	0.06	0.09	109
military	0.66	0.32	0.43	209
child_alone	0.00	0.00	0.00	0
water	0.76	0.64	0.69	415
food	0.80	0.69	0.74	731
shelter	0.77	0.53	0.63	576
clothing	0.73	0.45	0.56	97
money	0.59	0.29	0.39	140
missing_people	0.36	0.11	0.17	73
refugees	0.63	0.24	0.35	219
death	0.75	0.43	0.55	291
other_aid	0.54	0.14	0.22	855
infrastructure_related	0.42	0.10	0.16	427
transport	0.60	0.19	0.29	293

buildings	0.69	0.40	0.50	344
electricity	0.55	0.21	0.30	150
tools	0.50	0.03	0.06	34
hospitals	0.33	0.07	0.11	74
shops	0.17	0.04	0.06	26
aid_centers	0.23	0.10	0.14	72
other_infrastructure	0.34	0.08	0.13	294
weather_related	0.87	0.64	0.74	1818
floods	0.86	0.55	0.67	535
storm	0.75	0.49	0.59	612
fire	0.76	0.33	0.46	80
earthquake	0.89	0.75	0.82	627
cold	0.70	0.30	0.42	125
other_weather	0.42	0.13	0.20	337
direct_report	0.73	0.49	0.59	1273
avg / total	0.74	0.58	0.63	20593

/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWarning

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.

/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1137: UndefinedMetricWarning

Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples.

```
In [18]: # Print classification report on training data
print('\n',classification_report(y_train.values, y_prediction_train, target_names=y.col
```

	precision	recall	f1-score	support
related	0.82	0.95	0.88	14968
request	0.80	0.55	0.65	3355
offer	0.40	0.09	0.14	94
aid_related	0.77	0.61	0.68	8178
medical_help	0.66	0.30	0.42	1582
medical_products	0.71	0.34	0.46	1002
search_and_rescue	0.66	0.22	0.33	543
security	0.45	0.07	0.12	362
military	0.66	0.38	0.48	651
child_alone	0.00	0.00	0.00	0
water	0.78	0.68	0.73	1257
food	0.81	0.70	0.75	2192

shelter	0.81	0.57	0.66	1738
clothing	0.78	0.48	0.59	308
money	0.62	0.31	0.41	464
missing_people	0.78	0.22	0.34	225
refugees	0.66	0.28	0.39	656
death	0.80	0.48	0.60	903
other_aid	0.56	0.16	0.25	2591
infrastructure_related	0.54	0.11	0.18	1278
transport	0.75	0.24	0.36	908
buildings	0.72	0.43	0.53	989
electricity	0.64	0.25	0.36	382
tools	0.21	0.03	0.06	125
hospitals	0.45	0.15	0.23	209
shops	0.44	0.04	0.08	94
aid_centers	0.49	0.13	0.20	237
other_infrastructure	0.49	0.11	0.19	857
weather_related	0.87	0.66	0.75	5479
floods	0.87	0.58	0.70	1620
storm	0.77	0.52	0.62	1831
fire	0.62	0.31	0.41	202
earthquake	0.89	0.78	0.83	1828
cold	0.75	0.38	0.51	405
other_weather	0.59	0.22	0.32	1039
direct_report	0.76	0.49	0.60	3802
avg / total	0.76	0.60	0.65	62354

/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWarning

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.

/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1137: UndefinedMetricWarning

Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples.

1.0.8 9. Export your model as a pickle file

```
In [19]: m = pickle.dumps('classifier.pkl')
```

1.0.9 10. Use this notebook to complete train.py

Use the template file attached in the Resources folder to write a script that runs the steps above to create a database and export a model based on a new dataset specified by the user.

```
In [ ]:
```